A Public-Shift Softmax for Secure Inference of Bounded-Attention Transformers: Cost, Failure Modes, and the Price of the Guarantee

Pol Saumell Hill
Barcelona Supercomputing Center

Marc Casas Guix
Barcelona Supercomputing Center

Francesc Moll
Barcelona Supercomputing Center

Abstract

Secure two-party transformer inference computes softmax with a secure row maximum, a log-depth tree of comparisons. Recent systems drop it with an empirically chosen constant or clipping range, which is cheap but gives no guarantee on the inputs' range. We observe that architectures that bound their attention logits by construction (scalar-gain query-key normalization, attention-logit soft-capping) make a public constant shift *shift-exact for every input*: the shift cancels, and the exponential's argument lies in a known interval, so no comparison is needed. This is a guarantee about the numerical range, not about privacy; the secure exponential and reciprocal remain approximations. In two-party CrypTen the softmax takes a constant 37–38 rounds and, for a single softmax, communicates about 10× less than CrypTen's secure-max softmax and 13× less than SHAFT's. On BERT-base, measured two-machine wall time falls by 13% (128 tokens) and 24% (256 tokens) against SHAFT (mean of three runs, random weights); modelled LAN/WAN gains reach 35%/46%, and the advantage vanishes at 32 tokens. The guarantee has a price these costs do not include. The bound that is safe for CrypTen's 16-bit reciprocal ($g \leq 3$) costs converted BERT-base 2.9 to 14.8 GLUE points in plaintext under one untuned recipe; distillation and longer training recover at most 0.65 points on SST-2, while a ViT trained from scratch with the bound loses 0.3 ± 0.8 points (within noise, which with three seeds does not exclude a loss of one to two points). In two-party inference of trained models on 200 SST-2 sentences and 200 QNLI pairs, the shifted softmax has a lower median logit error than the baselines and is 5–7% faster at about 51 tokens (QNLI), but each protocol occasionally produces silent garbage logits (for the shifted softmax, about 1 sentence in 400 on each of SST-2 and QNLI). Released QK-normalized LLMs give vacuous bounds, so the bound must be designed in. On a 30-layer LLM, calibrated constants collapsed silently under CrypTen's reciprocal.

1 Introduction

Secure inference for transformers is dominated by non-linear functions. In a characterization of CrypTen-based private inference, softmax accounted for roughly half of the runtime of 12-layer language and vision transformers, and the computation of the maximum was singled out as a key culprit [11]. The reason is structural. The stable softmax subtracts $\max_j x_j$ so that every exponent is non-positive, but in secret-shared arithmetic a maximum is a tree of secure comparisons, each of which is a multi-round, non-polynomial protocol. Exponentiation and reciprocal, by contrast, admit smooth approximations.

A line of work therefore drops the maximum. NEXUS [14] uses a constant in its softmax, THOR [10] assumes inputs in a range taken from real data, and SHAFT [4] clips inputs to an empirically

tuned interval and then integrates an ODE; ARION [12] approximates the maximum by a predefined constant. These choices are *empirical*. The recent systematization of private transformer inference states the gap directly: whether a softmax protocol can be "non-interactive, efficient, and provably stable without distributional assumptions remains open" [5], and it suggests co-designing layers to obtain tighter bounds on the softmax inputs.

Idea. Some transformer architectures already bound their attention logits to a known interval, for stability reasons unrelated to secure computation. With query-key normalization [3] using a scalar gain g, every logit is $g \cos\theta \in [-g, g]$; with logit soft-capping [2] every logit lies in $(-c, c)$. For such a model the shift $m^* = g$ (or c) is a public constant that is valid for *all* inputs. It is exact because it cancels, and safe because the shifted logits lie in $[-2g, 0]$. No maximum, no comparison, and for the shift itself no calibration data and no guard (a trained model's other operations still need calibrated public scales, Section 3).

What we claim, and what we do not. A constant in place of the maximum is not new [14, 10, 12]; the claims here are narrower:

1. the bound makes the constant valid for every input and the range of the secure exponential and reciprocal known, and we state which conditions the fixed-point reciprocal needs; those conditions are measured for CrypTen, not proved (Section 3);
2. we measure the cost of the resulting protocol against CrypTen and SHAFT, from a single softmax up to full BERT-base and a GPT-2-shaped decoder, including the cost of QK-norm itself (Section 4);
3. we show, in 2-party inference of a 30-layer LLM, that calibrated constants fail silently and increasingly under distribution shift (Section 5);
4. we report the price: released models do not provide a usable bound, and the bound that is safe in fixed point costs real accuracy after conversion (Section 6).

We do not claim a free speed-up: the speed and the guarantee are bought with accuracy on converted checkpoints, and the speed and accuracy results come from different experiments (random-weight cost runs, plaintext fine-tuning). We tested distillation and longer fine-tuning (Section 6), which recover little; we have not evaluated QK-norm pretraining of a language model or a wider-window reciprocal, the levers most likely to reduce that cost.

2 Background and related work

Secure softmax. CrypTen [6] computes softmax as a secure maximum (a log-depth reduction), an exponential by $(1 + x/2^8)^{2^8}$ (eight squarings) and a Newton reciprocal. PUMA [1] and BumbleBee [9] also subtract the secure maximum and add a comparison per element to clip the exponential. SHAFT [4] removes the maximum by clipping each input to $[-4, 12]$ (two ReLUs per element) and integrating an ODE for 16 steps, for a constant 41 rounds, and states that convergence requires $\max(x) - \min(x) \le t$. NEXUS [14] and THOR [10] are homomorphic-encryption systems that use a constant or an assumed input range; NEXUS does not say how its constant is chosen, and the SoK [5] classifies constants of this kind, in these and other HE systems, as empirically determined. Many systems instead replace softmax with ReLU or quadratic forms and distill [7, 8]; those are approximate and need retraining, which differs from our setting, a shift-exact softmax on a bounded architecture.

Bounded attention. Query-key normalization [3] normalizes queries and keys to unit norm and scales by a learnable gain. Logit soft-capping $c\tanh(x/c)$ is used in Gemma 2 [2]. Both were introduced for training stability.

3 Method

3.1 Exactness and a priori bounds

For any scalar m^*, $\mathrm{softmax}(x)_i = e^{x_i - m^*} / \sum_j e^{x_j - m^*}$, since m^* cancels. The maximum is needed only to keep $x_j - m^* \le 0$. Hence any public $m^* \ge \max_j x_j$ is exact and safe, and the question is whether $\max_j x_j$ can be bounded *before* seeing the input.

Proposition 1 (QK-norm)**.** *Let $\hat{q}_i = q_i/\|q_i\|_2$, $\hat{k}_j = k_j/\|k_j\|_2$ and $x_{ij} = g\,\hat{q}_i \cdot \hat{k}_j$ with a scalar $g > 0$. Then $x_{ij} \in [-g, g]$ for every input, and $m^* = g$ gives shifted logits in $[-2g, 0]$.*

Proof. Cauchy–Schwarz: $|\hat{q}_i \cdot \hat{k}_j| \le \|\hat{q}_i\|\|\hat{k}_j\| = 1$. □

Proposition 2 (Soft-capping)**.** *If $x = c\tanh(r/c)$ then $|x| < c$ for every raw score r, so $m^* = c$ is valid before training.*

With a row of k visible keys the shifted exponentials lie in $[e^{-2g}, 1]$ and the row sum lies in $[k\,e^{-2g}, k]$, a known interval. Vanilla attention has no such bound: $|x| \le \|W_q x\|\|W_k x\|/\sqrt{d}$ depends on weight norms that can grow during training [13].

What "exact" means here. The shift is exact because it cancels; the protocol's output is not exact. The secure exponential $(1 + x/2^8)^{2^8}$ and the Newton reciprocal are approximations (the exponential's log error is about -14% at $x = -8$), and the worst-case row errors we measured are 0.7–9.5% (Section 3.4). We say "shift-exact" for the first property. "Provable" below means a guaranteed range for the exponential's argument; it is not a security or end-to-end accuracy guarantee.

3.2 Threat model and what is public

We consider semi-honest two-party computation with additive secret sharing, as in CrypTen. We make no new security argument: the protocol composes existing CrypTen/SHAFT primitives (exponential, reciprocal, linear operations), and the shift is data-independent. What is public must be stated, because it differs between our experiments: the architecture (that QK-norm is used), the sequence length N, and an upper cap $g_{\max}$ used to choose the reciprocal window and the public rescale c. The microbenchmark and the ViT experiment use a *public* g (the shift $m^* = g$ is a public constant); the BERT and GPT-2 cost runs use the *shifted* form, where the logits are formed as $g\cos\theta - g$ so that the shift is folded into a linear operation on shared values and the learned g is not revealed. If the model weights are private, a public per-layer g leaks a trained parameter, so the shifted form is the appropriate one; but the rescale c and Eq. (1) are then chosen from the public cap $g_{\max}$, not from the learned g. The LLM experiments use public weights and secret activations; the cost runs use random encrypted weights. The public quantities are therefore: the architecture, the sequence length N, the cap $g_{\max}$, the rescale c, and the per-layer inverse-square-root scales c^2 of Section 3. The last are calibrated on training data and are model-dependent (up to about 200 for late LayerNorms), so they reveal coarse activation statistics of the model; we did not analyse this leakage.

3.3 Secure protocol

Given the shared logits, the protocol is: subtract m^* (a public constant, or a shared parameter, in which case g is not revealed because subtracting a shared weight is a free linear operation), apply the secure exponential, sum, and apply the secure reciprocal. No comparison appears. The cost is 37–38 rounds per call, independent of the row length. QK-norm adds its own secure operations, an inverse square root per token and head, which cost 300 extra rounds for BERT-base; we include them in every number below.

Public scales for the inverse square roots. Running a trained BERT end to end exposed a second range problem that random-weight cost runs hide. CrypTen's Newton inverse square root, with its fixed initial guess, diverges for inputs above roughly 75. In the trained SST-2 models the sum of squares of a query or key head vector reaches several hundred, and the variance of a LayerNorm input row reaches about 1,500–1,800 in layers 9–10 (activation outliers); without a remedy QK-norm's normalization failed at layer 2, and the vanilla baseline with CrypTen's own softmax failed at layer 9. We divide each such input by a public per-layer constant c^2 before the inverse square root and fold $1/c$ back in ($1/\sqrt{x} = \text{inv_sqrt}(x/c^2)/c$), which adds no rounds. The constants are calibrated on 128 *training* sentences (never the evaluation ones) as $c^2 = \max(1,\, 3 \cdot \max_{\text{calib}}/25)$, giving $c^2 = 1$ for most layers and up to about 200 for the late LayerNorms. We apply them identically to every variant, baselines included. The shift itself needs no calibration, but the surrounding model does.

3.4 Fixed-point conditions: what "provable" requires

The bound makes the numerators representable but not automatically the reciprocal safe, and this is the condition that matters in practice. We measured CrypTen's Newton reciprocal (16-bit fixed point, 10 iterations, 2-party, 64 repetitions per point) and found it within 1% only for $\ln s \in [-6.5, +3.0]$; the error is 4.3% at +4.0, 30% at +5.0, and the output is garbage for $\ln s \geq 5.5$ (a row sum ≥ 256); at the low end it is 2.6% at -8 and 36% at -9 (under-converged). A public rescale $1/s = (1/(s/c))/c$ costs two public multiplications and no communication, and shifts the window. With $c = \max(1, N/64)$ and $k_{\min}$ the fewest visible keys in any row, we find the lower end is the binding constraint:

$$2g + \ln c \;\leq\; 6.75 + \ln k_{\min}. \tag{1}$$

At $k_{\min} = 1$ (a causal decoder) and $N \leq 128$ this gives $g \leq 3$; $g = 5.5$ is not safe, because a row whose only visible key sits at $-2g$ has $\ln s = -11$, far outside the window. We stress that (1) is derived from measurements of this specific reciprocal, not proved, and holds for the tested range ($N \leq 512$, six row kinds, $g \in \{3, 5.5\}$).

4 Cost evaluation

Setup. All protocols run in 2-party CrypTen using SHAFT's fork (torch 2.0.1), 16-bit fixed point, on MareNostrum 5 general-purpose nodes (Xeon 8480+). Cost is weight-independent, so BERT/GPT-2 costs use random encrypted weights; rounds and bytes are exact and deterministic. Three softmaxes: CrypTen's secure max; SHAFT's clip + ODE ($t = 16$); and ours (QK-norm + shifted softmax). Network time uses SHAFT's model (compute + 8·bytes/bandwidth + rounds·latency; LAN 3 Gbps/4 ms, WAN 400 Mbps/40 ms); we also measure wall-clock on two machines.

Table 1: BERT-base encoder, 16-bit, 2 parties, random weights (N = tokens). Wall time is the mean of three runs, measured on two machines; LAN/WAN are modelled.

N	Softmax	Rounds	GB	Wall (s)	LAN (s)	WAN (s)
32	CrypTen max	2208	2.78	–	32.8	160.5
	SHAFT	1440	2.85	–	29.5	130.7
	Ours	1704	2.61	–	29.5	136.1
128	CrypTen max	2340	9.70	54.9	90.0	342.5
	SHAFT	1440	10.80	55.3	90.3	329.3
	Ours	1704	6.68	**47.9**	**72.8**	**249.8**
256	CrypTen max	2460	24.16	116.8	192.1	699.5
	SHAFT	1440	29.52	124.0	209.5	773.1
	Ours	1704	12.83	**94.1**	**135.6**	**419.3**

Notes. Wall: two Xeon 8480+ nodes (23 Gbps link, 56 threads per party), N=128 and 256 only. Rounds, GB, LAN and WAN come from the single-node run; LAN/WAN apply SHAFT's time model to that run's compute time, so they mix two runs with the Wall column (their totals differ by up to about 1 s). The Wall column is the mean of three runs on the same two nodes (standard deviation at most 0.3 s); rounds and bytes are deterministic; LAN and WAN come from one run. CrypTen and SHAFT run unmodified, without the public rescale given to ours, so the CrypTen cells at N=256 are not like-for-like.

4.1 Single softmax

Per 16-row batch, $n = 1024$ (inputs in $[-g, g]$): CrypTen 137 rounds, 22.8 MB, 185 ms; SHAFT 41 rounds, 32.0 MB, 129 ms; ours 37 rounds, 2.37 MB, 41 ms. Ours is 9.6× and 13.5× smaller in communication and 4.5×/3.1× faster. CrypTen's secure-max softmax also breaks at long rows: at $n = 1024$, flat rows give errors of 10^7–4×10^9 at both precisions, while ours stays at 10^{-4}. Removing the maximum removes that failure mode too. SHAFT's ODE is about 11× less accurate than ours on peaked rows (max error 0.17 vs 0.015).

4.2 BERT-base and a causal decoder

Table 1 gives BERT-base (12 layers, hidden 768) encoder cost. Softmax shrinks from 38–55% of the bytes to a few percent (Figure 1); GELU becomes the largest cost. QK-norm costs 300 extra rounds (total 1704 vs SHAFT's 1440, +18%) but only 0.1–0.2 GB.

Against SHAFT, the measured two-machine wall time falls by 13.4% at N=128 and 24.1% at N=256; against CrypTen by 12.9% and 19.5% (means of three runs). End-to-end bytes fall by 31% and 38% at 128 tokens and by 47% and 57% at 256 tokens against CrypTen and SHAFT (the 10× and 13× figures belong to the single softmax). The wall-time gap against SHAFT is mostly SHAFT's 16-step ODE: at N=128 the softmax takes 9.5 s for SHAFT, 8.9 s for CrypTen's secure max and 1.3 s for ours, so against CrypTen the gain is the smaller one. The three runs agree to within 0.3 s (standard deviation), and the two-node and single-node totals agree within 1.5%. The modelled LAN/WAN reductions against SHAFT are 19.4%/24.1% at 128 and 35.3%/45.8% at 256. With the embedding, pooler and classifier included (N=128) they are 17.4%/22.7%. A GPT-2-shaped causal decoder (prefill, blocks only, public mask sent to −30) behaves the same: −19.3% (LAN) and −24.0% (WAN) at 128, and −35.0% and −45.7% at 256. Under QK-norm the shifted logits lie in $[-2g, 0]$, so the −30 mask rounds to exactly zero.

Caveats on these numbers.

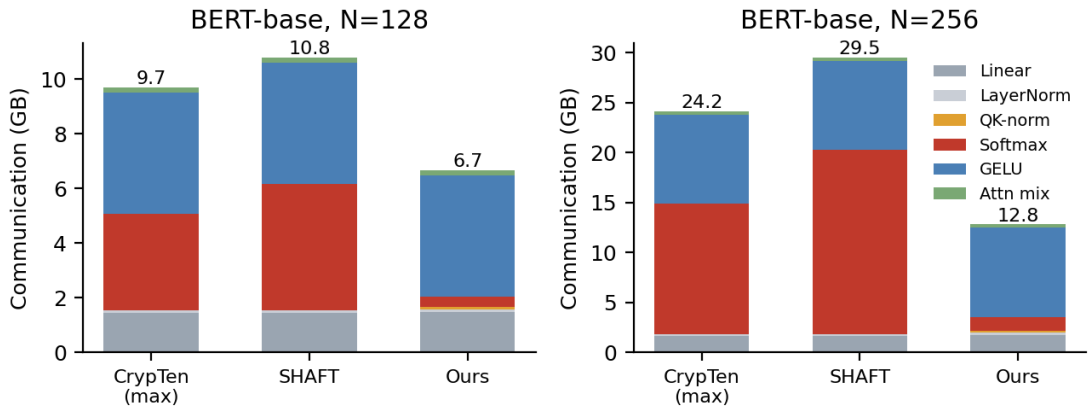


Figure 1: Communication by component, BERT-base, two-node run. Softmax dominates the baselines and is small in ours; QK-norm adds a small slice.

1. At N=32 the extra QK-norm rounds cancel the byte savings (LAN +0.2%, WAN +4.2% vs SHAFT): there is no gain for short inputs.
2. The two-node link is fast (23 Gbps, 0.16 ms) and communication is only 3–5% of wall time there, so the measured advantage is a compute advantage. The LAN and WAN rows remain *modelled*; the model over-predicts measured communication by 1.35–1.65× on this link (with the byte term scaled by 0.6 the reductions at 128 tokens are 16.3%/19.0%, same ranking). A real WAN was not run.
3. Raising the Newton iterations to 20 for a wider reciprocal window adds 240 rounds and moves the 128-token gain to −18.1% (LAN) and −21.2% (WAN).
4. A fair baseline would also receive the public rescale, which removes CrypTen's own overflow at $N \geq 256$. We did not run CrypTen or SHAFT with it, so the N=256 comparison against CrypTen is not like-for-like, and SHAFT's clip range was not tuned to bounded logits. We also did not run a constant shift *without* QK-norm: its softmax cost would equal ours minus QK-norm's 300 rounds, but it carries no bound, so we could not run it safely and report only the expectation.
5. Table 1 is measured with random weights. Rounds and bytes do not depend on the values, but the fixed-point behaviour does, and trained weights needed the public scales of Section 3. Section 6 reports a trained model at about 26 tokens, where there is no speed-up.
6. Our SHAFT total at N=128 with the embedding (11.25 GB) differs from SHAFT's paper (10.46 GB), plausibly because of our one-hot embedding and a CrypTen-level reimplementation of the encoder (SHAFT's ONNX path did not import the QK-norm graph).
7. These costs do not depend on g. The accuracy of models that satisfy Eq. (1) is in Section 6.

4.3 Worst-case rows

We ran 16 rows each of six kinds (all logits at $-2g$, all at 0, one at 0, one visible key at $-2g$ as in a causal first token, uniform over $[-2g, 0]$, and trained-like cosines) for $g \in \{3, 5.5\}$, $N \in \{32, \ldots, 512\}$. With $g \leq 3$ and $c = \max(1, N/64)$ no row produced garbage up to N=512; the worst row-sum error is $\leq 0.7\%$ to 128 tokens and 1.45% at 256, versus 0.1–2.4% for CrypTen's own softmax, which

Table 2: Secure SmolLM2-135M, 16-bit, 9 windows (v2). Collapsed = non-finite, top-1 < 0.5 or perplexity $\times 2$.

Softmax	Collapsed W/code/Catalan	Mean top-1 (W/code/ca)	LAN / WAN (s)
CrypTen secure max	0/3, 0/3, 0/3	.98 / .99 / .98	55.5 / 337.3
Constant per layer	3/3, 3/3, 3/3	.02 / .00 / .01	38.0 / 212.0
Constant per head	3/3, 3/3, 3/3	.01 / .02 / .04	38.0 / 212.0
Constant per head×position	3/3, 3/3, 2/3	.04 / .27 / .53	38.0 / 212.0

produces garbage on uniform rows at $N \geq 256$ because its row sums reach N. At 512 our worst case is 9.5% (the one-key causal row loses resolution when divided by $c = 8$). With $g = 5.5$, the one-key row ($\ln s = -11$, outside the window; 87% error in the fine sweep) produced garbage at every N tested, as Eq. (1) predicts. Calls take 38 rounds against 106–137.

5 Calibrated constants on an unbounded LLM under CrypTen's reciprocal

Since a calibrated constant needs no architecture change, the natural alternative is to profile one. We ran SmolLM2-135M (all 30 decoder layers, weights public, activations secret, embedding lookup and LM head in the clear) in 2-party CrypTen at 16 bits on 64-token windows. The plaintext reimplementation matches Hugging Face (KL $5\text{–}9 \times 10^{-6}$, identical perplexity). Constants were calibrated on 128 Wikitext windows (+1 nat) at layer, head, and head×position granularity.

Plaintext statistics. 43–46% of rows underflow at 16 bits with a per-layer constant, and 4–5% with a per-head constant, even on the calibration domain. At the finest granularity the share of windows containing a row the reciprocal cannot handle is 3.1% on held-out Wikitext, 0–1.6% on Spanish/Catalan, 6.2% on math, 17.2% on chat and 31–39% on code, with the per-row excess over the constant reaching +12.5 nats on code (+7.7 on Wikitext). Separately, a layer-wide constant profiled on Wikitext is exceeded by 9.2–9.4 nats on code for both SmolLM2 and TinyLlama and by 4–5.6 nats on Catalan and Spanish (plaintext attention); FP32 absorbs this, 16-bit secure arithmetic does not.

Secure runs. With CrypTen's secure maximum, 9/9 windows are fine (top-1 agreement 0.98–1.00 with plaintext, perplexity within about 1–2%). With a calibrated constant every granularity collapses (Table 2; LAN/WAN are modelled from per-window times aggregated as in our run logs, and recomputing from the raw CSV mean gives values 1–2 s higher): per-layer and per-head constants fail on all windows; the finest per-row constant fails on 8 of 9, including in-distribution windows; 20 Newton iterations make it worse (9/9 collapsed). The mechanism is that a handful of rows per layer (0.2–0.3% underflow, 2–3% short of the reciprocal's range) produce a reciprocal that overflows the 2-party truncation, and one such row per layer is enough across 30 layers. With a public floor e^{-6} on the row sum (job v3, 32 windows, four domains), outcomes are *garbage* on 3/8 held-out Wikitext windows and 6/8 HumanEval windows. The failure rate rises with how far the data exceed the constant (Figure 2): 27% (excess < 3 nats), 46% (3–6), 88% (≥ 6). Outputs stay finite, so a non-finite guard would not fire.

Three scope notes. (1) This is a statement about CrypTen's 16-bit reciprocal: the collapse is a cliff caused by a few rows whose sum leaves the reciprocal's working range, and a range-extended

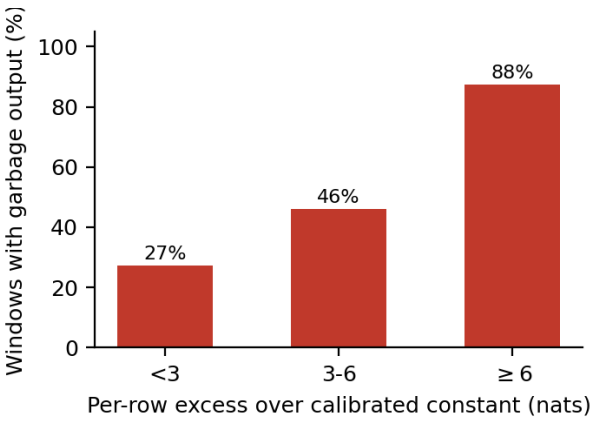


Figure 2: Fraction of windows with garbage output vs per-row excess over the calibrated constant (floor e^{-6}, 32 windows). Few windows per bin (8–13); read the trend, not individual rates.

reciprocal, more precision, or a clamp of the row sum might rescue a calibrated constant; we did not try these, so the section shows that calibrated constants are fragile under this reciprocal, not that they cannot work. (2) We ran 8–9 windows per domain and one model, with bins of 8–13 windows built from four domains and no confidence intervals, so read the trend, not individual rates. (3) We could not separate this failure from model type: a test of input-level distribution shift on pre-LN ViTs never pushed attention past a clean-calibrated constant (0 of 266,240 rows), so we make the claim only for language models.

6 The price of the guarantee

Released models give no bound. Computed from the released weights, the logit bounds of per-dimension QK-norm models are enormous: 184–12,897 nats for Qwen3-0.6B, 183–5,361 for Qwen3-1.7B, and 995–4,267 for OLMo-2-1B; Gemma 2's soft-cap $c = 50$ gives a 100-nat width. None fits a 16-bit range (≈ 11 nats). A usable bound must be designed in: a scalar gain with a small cap, trained or fine-tuned in.

Vision. A CIFAR-10 ViT trained from scratch with scalar QK-norm reaches 81.74% against 81.24% vanilla, logits in $[-4.29, 4.38] \subset [-4.48, 4.48]$, so $m^* = g$ is exact without profiling. In 2-party inference of that model (200 images, one run, public g) the shifted softmax gives 99.0% agreement with plaintext (CrypTen max 98.5%, SHAFT 99.0%) with the least communication of five variants (94.3 vs 177.8–215.9 MB per image). At 65 tokens the extra QK-norm rounds cancel the byte savings (WAN +3% vs vanilla+SHAFT).

Converted BERT-base. We convert BERT-base to scalar-gain QK-norm and fine-tune it for 3 epochs (lr 2×10^{-5}, gains at 10^{-3}, one recipe for all tasks, not tuned). Table 3 shows the result: the bound that Eq. (1) requires ($g \leq 3$) costs 2.9–14.8 points under this single untuned recipe, which may overstate the cost of the bound itself (the gains sit at the cap; controls with a longer, distilled and higher-lr recipe, and a from-scratch ViT control, follow below); even $g \leq 5.5$, which is not provable for single-key rows, costs 1.0–4.8. SST-2 is the mildest task. The gains sit at the cap on every task, so the bound binds. The mechanism is that g caps attention sharpness: the largest weight a single key can receive among N is $e^{2g}/(e^{2g} + N - 1)$, i.e. 0.76 at $g = 3$, $N = 128$ (0.998 at

Table 3: GLUE dev, BERT-base converted from the vanilla checkpoint (mean ± sd over seeds: 5 for RTE, MRPC and CoLA; 3 for QNLI, MNLI and SST-2). Metric: accuracy; MRPC F1; CoLA Matthews.

Task	Vanilla	$g \leq 5.5$	$g \leq 3.0$ (fixed-point safe)
SST-2	92.89±0.30	91.93±0.29	89.95±0.63
QNLI	91.48±0.20	89.33±1.48	86.36±0.90
MNLI-m	84.38±0.17	83.16±0.30	80.58±0.12
MRPC (F1)	88.57±1.25	85.10±0.80	83.14±1.32
RTE	62.31±0.94	57.47±2.47	55.02±2.48
CoLA (MCC)	57.28±1.43	54.42±1.02	42.53±3.22

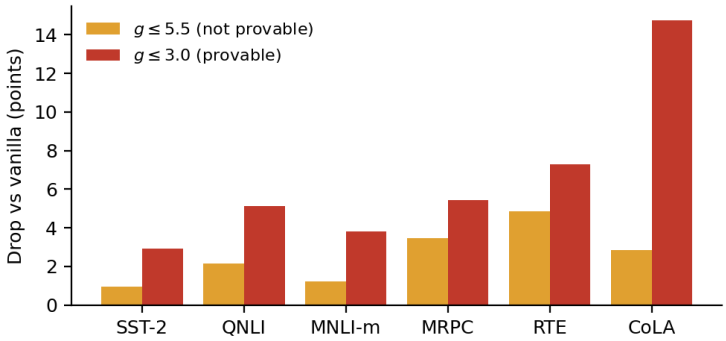


Figure 3: Accuracy drop of QK-norm-converted BERT-base against vanilla.

$g = 5.5$), which may hurt tasks needing token-level attention (CoLA, RTE, MRPC); the controls below suggest that this explains only part of the loss.

How much of the cost is the bound, and how much the recipe? We ran two controls, both with the cap $g \leq 3$ enforced (learned gains 2.85–2.98, logit width under 6 nats). *BERT-base on SST-2, recipe variants* (Table 4, 3 seeds each): distilling from the vanilla fine-tuned model, a higher learning rate, or 10 epochs alone leave accuracy within noise of the baseline recipe (89.83–90.06 against 89.95); distillation with 10 epochs gives 90.60 ± 0.41, which narrows the gap to vanilla (92.89) from 2.9 to 2.3 points. That 0.65-point gain is about one standard deviation with three seeds, so it is suggestive only, and the controls were scored on the dev set that Table 3 also uses, so they are a sensitivity analysis, not a tuned result. *ViT trained from scratch with* $g \leq 3$ (CIFAR-10, 100 epochs, 3 seeds): vanilla 81.32 ± 0.19 against QK-norm 81.01 ± 0.74 (QK-norm seeds 81.50, 80.15, 81.39; vanilla 81.24, 81.18, 81.53). The learned gain is 2.84, the logits lie in $[-2.79, 2.79]$, and the shift error is about 10^{-7}. The 0.3-point gap is within seed noise. Taken together, the cap alone does not account for the large conversion losses: a model trained with the bound from the start loses nothing measurable here, and much of the BERT loss comes from converting a pre-trained model in a few epochs. This is thin evidence (a small ViT, one dataset, and no pre-training with the bound), and it does not show that a bounded BERT pre-trained from scratch would match vanilla.

Secure inference of the converted BERT. We ran the SST-2 models end to end in two-party CrypTen (16-bit): the $g \leq 3$ QK-norm model with the shifted softmax, and the vanilla model with CrypTen's secure max and with SHAFT's ODE, on the same 200 validation sentences (a

Table 4: BERT-base SST-2 with $g \leq 3$: recipe variants (dev accuracy %, 3 seeds). KD = distillation from the vanilla fine-tuned model.

Recipe	Epochs	lr	Mean ± sd
Baseline	3	2×10^{-5}	89.95±0.63
KD	3	2×10^{-5}	89.83±0.47
Higher lr	3	5×10^{-5}	90.06±0.52
Longer	10	2×10^{-5}	90.02±0.69
KD, longer	10	2×10^{-5}	90.60±0.41

Table 5: Secure two-party inference of fine-tuned BERT-base on 200 SST-2 validation sentences (mean 25.9 tokens), 16-bit, one seed; same public scales for every row. Plain/Secure = accuracy (%); err = max|secure − plaintext logit| per sentence; garb. = sentences with err > 1. Costs are means per sentence.

Model, softmax	Plain	Secure	Agree	Med. err	Garb.	Rounds / GB	s
QK-norm $g \leq 3$, shifted (ours)	91.0	91.0	200/200	0.023	1	1790 / 2.40	15.5
Vanilla, CrypTen max	92.5	92.0	199/200	0.033	1	2211 / 2.53	16.5
Vanilla, SHAFT ODE	92.5	92.0	199/200	0.036	0	1490 / 2.57	15.7

seeded random subset; mean 25.9 tokens, no padding, so every row has all keys visible). Weights are secret-shared by party 0; the embedding output of each sentence is computed by the client (party 1) and encrypted; the encoder, pooler and classifier are secure. The $g \leq 3$ model is the seed-0 model from Table 3 (90.25% on the full dev set). Table 5 gives the results. The converted model's secure predictions agree with its plaintext predictions on 200 of 200 sentences, and the secure accuracy equals its plaintext accuracy (91.0%), 1.0 point below the secure vanilla model (92.0%), as in plaintext on this subset. **This agreement overstates the numerical quality:** on one sentence (41 tokens) the secure logits were garbage (-8×10^8 and -2×10^8 against plaintext -3.0 and $+3.1$), and the class was right only because the garbage values happened to be ordered the right way round. The output was finite, so a non-finite guard would not fire. CrypTen's secure max failed the same way once (26 tokens), and there it flipped the prediction; SHAFT's ODE produced no garbage logits (maximum error 0.68) but disagreed with plaintext once. Rerunning each of the two failing sentences alone, three times with fresh randomness, gave no failure (6 of 6 runs with maximum error 0.02–0.05), so the failures are not determined by the input; they depend on the shares or the run state. A full second pass of the $g \leq 3$ model over the same 200 sentences with fresh randomness had no sentence with error above 1, including the one that failed before, and its logits differ from the first pass by a median of 0.002 (maximum 0.019). The failure is therefore a rare event driven by the randomness of the shares, about 1 in 400 sentence runs for this model on SST-2. We did not identify the cause (probabilistic truncation in CrypTen is a candidate we did not test), and with such a low rate we cannot rank the softmaxes. At this sequence length the converted model is not meaningfully faster (15.5 s against 16.5 s for the secure max and 15.7 s for the ODE per sentence) and uses more rounds than the ODE, consistent with the absence of a gain at 32 tokens in Table 1.

Longer inputs: QNLI. SST-2 sentences are short, so we repeated the experiment on QNLI (question and sentence pairs; 200 seeded validation pairs, mean 51.3 tokens, up to 128), with the same recipe and public scales; the main $g \leq 3$ model is the seed-0 model (87.0% on the full dev set, against 91.3% for vanilla); a second seed is reported below. Table 6 gives the results. The shifted

Table 6: Secure two-party inference of fine-tuned BERT-base on 200 QNLI validation pairs (mean 51.3 tokens), 16-bit, seed-0 converted model; same public scales for every row. Err > 1 / > 10 = sentences whose maximum logit error exceeds 1 / 10. Costs are means per sentence.

Model, softmax	Plain	Secure	Agree	Med. err	Err>1	Err>10	Rounds / GB	s
QK-norm $g \leq 3$, shifted (ours)	86.5	87.0	199/200	0.074	1	1	1790 / 3.45	23.9
Vanilla, CrypTen max	91.0	91.0	198/200	0.395	8	0	2305 / 3.95	25.9
Vanilla, SHAFT ODE	91.0	91.0	196/200	0.448	19	1	1490 / 4.14	25.5

softmax is faster than both baselines on every sentence pair compared: by 7.4% against CrypTen's secure max and 5.5% against SHAFT's ODE (per-sentence paired differences; 95% intervals ± 0.2 and ± 0.3 points, which reflect sentence-to-sentence variation only, since the three runs were separate jobs on different nodes of the same type), with 11% and 15% fewer bytes. The median logit error is five times smaller (0.074 against 0.40 and 0.45), and 1 sentence has an error above 1 against 8 and 19. The catastrophic failure recurs: one sentence (72 tokens) produced garbage logits ($\sim 10^5$ in error) under the shifted softmax, and one under SHAFT's ODE (a different sentence). Rerun alone three times each, neither reproduced (errors 0.08 and 0.23), as on SST-2, so 12 of 12 isolated reruns of the four failing sentences were clean; CrypTen's max produced none here, although it had one in the SST-2 run. Secure and plaintext predictions agree on 199, 198 and 196 of 200 sentences; the flipped predictions of the $g \leq 3$ model fall on a sentence whose plaintext logits differ by 0.014, so they are near-ties. The secure accuracy of the converted model stays 4 points below the secure vanilla model, as in plaintext. A second QNLI seed of the converted model (full dev accuracy 86.75%, against 87.00%) behaves similarly and has no catastrophic failure: secure accuracy 88.5% (plaintext 88.0%) on the same 200 pairs, agreement 199/200, median error 0.035, no sentence with error above 1, the same 1790 rounds and 3.45 GB, and 6.6% and 4.7% faster than CrypTen's max and SHAFT's ODE (it is faster than the ODE on 194 of 200 pairs). The catastrophic failure therefore occurred for one of the two QNLI models (1 sentence in 400 for the shifted softmax on QNLI); the two seeds differ in their weights, so these are not two draws of the same experiment. The speed gain at 51 tokens is much smaller than the 13–24% measured at 128–256 tokens with random weights, as the cost model predicts.

7 Discussion and limitations

Where this leaves the trade-off. The speed result (Section 4) and the accuracy result (Section 6) are measured on the same protocol but not on the same operating point as a deployment would choose: the cost does not depend on g, while the guarantee does. The honest summary is that, for a model that is *already* bounded (or can be retrained so), a shift-exact, comparison-free softmax is a clear communication saving; for a converted model it is a speed/accuracy trade of several points on harder tasks. Three levers could change the trade-off and were not tested: a reciprocal with a wider working window (the binding constraint is the reciprocal, not the model), QK-norm pretraining (distillation and longer fine-tuning were tested and recover at most 0.65 points on SST-2), and a softmax with a constant in the denominator ("softmax-one") or a public floor on the row sum. The last changes the function (it is no longer the exact softmax) unless it is a floor applied only for numerical safety, and a floor of e^{-6} did not rescue calibrated constants on the LLM.

Limitations.

1. Fixed-point conditions come from measuring one reciprocal implementation (CrypTen's), not from a proof.
2. LAN/WAN numbers use SHAFT's time model; wall-clock is on a fast link.
3. The large cost tables use random weights. Secure accuracy of the converted BERT was measured on two tasks (SST-2, QNLI), 200 sentences each (about 26 and 51 tokens), one seed, and not on the other GLUE tasks.
4. Decoder results are prefill cost only and no QK-norm decoder was trained. Decode-phase cost with a KV cache, where a secure maximum over one row costs about $\log_2 t$ rounds, is not measured, and the advantage there is likely smaller.
5. The calibrated-constant failure is for an unbounded LLM at 16 bits in CrypTen, from few windows.
6. We evaluate semi-honest 2-party computation in one framework; the argument that the shift can be public or shared composes existing protocols and is not a new security proof.
7. Weights are public in the LLM experiments; weight privacy changes cost, not these numerics.
8. The $g \leq 3$ ViT checkpoints were evaluated in plaintext only; the 2-party ViT run uses the earlier unbounded-gain model.
9. Only the QK-norm case is implemented and costed; soft-capping needs a secure tanh that we did not implement.
10. The 16-bit ring is CrypTen's choice; most current 2PC systems use larger rings, and the 11-nat window, hence $g \leq 3$, may be an artefact of that precision.
11. Wall times are three-run means on one pair of nodes (the repeats share hardware, so node-to-node variation is not captured), and LAN/WAN are modelled from one run; the first ViT result (81.7% vs 81.2%) is one run, and the $g \leq 3$ ViT control uses 3 seeds on a small model and one dataset.

8 Conclusion

A bounded attention architecture turns the softmax's numerical-safety shift into a public constant that is shift-exact for every input, eliminating the secure maximum. On BERT-base this saves 13–24% of two-machine wall time against SHAFT at 128–256 tokens (three-run means, random weights), and nothing at 32. The guarantee is not free: calibrated constants are not a substitute (they fail silently on an LLM), released models give no usable bound, and the bound that is safe in fixed point costs real accuracy after conversion. A converted $g \leq 3$ BERT does run end to end in two-party inference, with secure accuracy equal to plaintext on 200 SST-2 sentences and 200 QNLI pairs, a five- to six-fold lower median logit error than the baselines and a 5–7% speed-up at about 51 tokens, but with an occasional silently corrupted output (about one sentence in 400 on each of SST-2 and QNLI, not reproduced when rerun alone) and no speed-up at 26 tokens. A $g \leq 3$ ViT trained from scratch matches vanilla within seed noise, which weakly suggests that some of the conversion loss comes from the recipe rather than from the bound; the evidence (one small vision model, three seeds) cannot establish this for language tasks that need sharp attention. The open question is whether a language model pre-trained with the bound can have the speed, the guarantee and the accuracy together.

Availability. The benchmark and training scripts, job scripts and per-sentence result files are available at `https://github.com/polsahi/soft-no-max`; the fine-tuned checkpoints (about 4 GB) are available from the authors on request.

Acknowledgments. The experiments ran on MareNostrum 5 at the Barcelona Supercomputing Center.

References

[1] Ye Dong, Wen-jie Lu, Yancheng Zheng, Haoqi Wu, Derun Zhao, Jin Tan, Zhicong Huang, Cheng Hong, Tao Wei, and Wenguang Chen. PUMA: Secure inference of LLaMA-7b in five minutes, 2023. arXiv:2307.12533.

[2] Gemma Team, Google DeepMind. Gemma 2: Improving open language models at a practical size, 2024. arXiv:2408.00118.

[3] Alex Henry, Prudhvi Raj Dachapally, Shubham Shantaram Pawar, and Yuxuan Chen. Query-key normalization for transformers. In *Findings of EMNLP*, 2020.

[4] Andes Y. L. Kei and Sherman S. M. Chow. SHAFT: Secure, handy, accurate, and fast transformer inference. In *Network and Distributed System Security Symposium (NDSS)*, 2025.

[5] Andes Y. L. Kei and Sherman S. M. Chow. SoK: Private transformer inference across systems, models, and cryptography, 2026. IACR ePrint 2026/2005, preprint received September 2026.

[6] Brian Knott, Shobha Venkataraman, Awni Hannun, Shubho Sengupta, Mark Ibrahim, and Laurens van der Maaten. CrypTen: Secure multi-party computation meets machine learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.

[7] Dacheng Li, Hongyi Wang, Rulin Shao, Han Guo, Eric Xing, and Hao Zhang. MPCFormer: Fast, performant and private transformer inference with MPC. In *International Conference on Learning Representations (ICLR)*, 2023.

[8] Zhengyi Li, Kang Yang, Jin Tan, Wen-jie Lu, Haoqi Wu, Xiao Wang, Yu Yu, Derun Zhao, Yancheng Zheng, Minyi Guo, and Jingwen Leng. Nimbus: Secure and efficient two-party inference for transformers. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

[9] Wen-jie Lu, Zhicong Huang, Zhen Gu, Jingyu Li, Jian Liu, Cheng Hong, Kui Ren, Tao Wei, and WenGuang Chen. BumbleBee: Secure two-party inference framework for large transformers. In *Network and Distributed System Security Symposium (NDSS)*, 2025.

[10] Jungho Moon, Dongwoo Yoo, Xiaoqian Jiang, and Miran Kim. THOR: Secure transformer inference with homomorphic encryption. In *ACM Conference on Computer and Communications Security (CCS)*, 2025. IACR ePrint 2024/1881.

[11] Yongqin Wang, G. Edward Suh, Wenjie Xiong, Benjamin Lefaudeux, Brian Knott, Murali Annavaram, and Hsien-Hsin S. Lee. Characterization of MPC-based private inference for transformer-based models. In *IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, 2022.

[12] Linhan Yang, Jingwei Chen, Wangchen Dai, Shuai Wang, Wenyuan Wu, and Yong Feng. ARION: Attention-optimized transformer inference on encrypted data, 2025. IACR ePrint 2025/2271.

[13] Shuangfei Zhai, Tatiana Likhomanenko, Etai Littwin, Dan Busbridge, Jason Ramapuram, Yizhe Zhang, Jiatao Gu, and Joshua M. Susskind. Stabilizing transformer training by preventing attention entropy collapse. In *International Conference on Machine Learning (ICML)*, 2023.

[14] Jiawen Zhang, Xinpeng Yang, Lipeng He, Kejia Chen, Wen-jie Lu, Yinghao Wang, Xiaoyang Hou, Jian Liu, Kui Ren, and Xiaohu Yang. Secure transformer inference made non-interactive. In *Network and Distributed System Security Symposium (NDSS)*, 2025. IACR ePrint 2024/136.